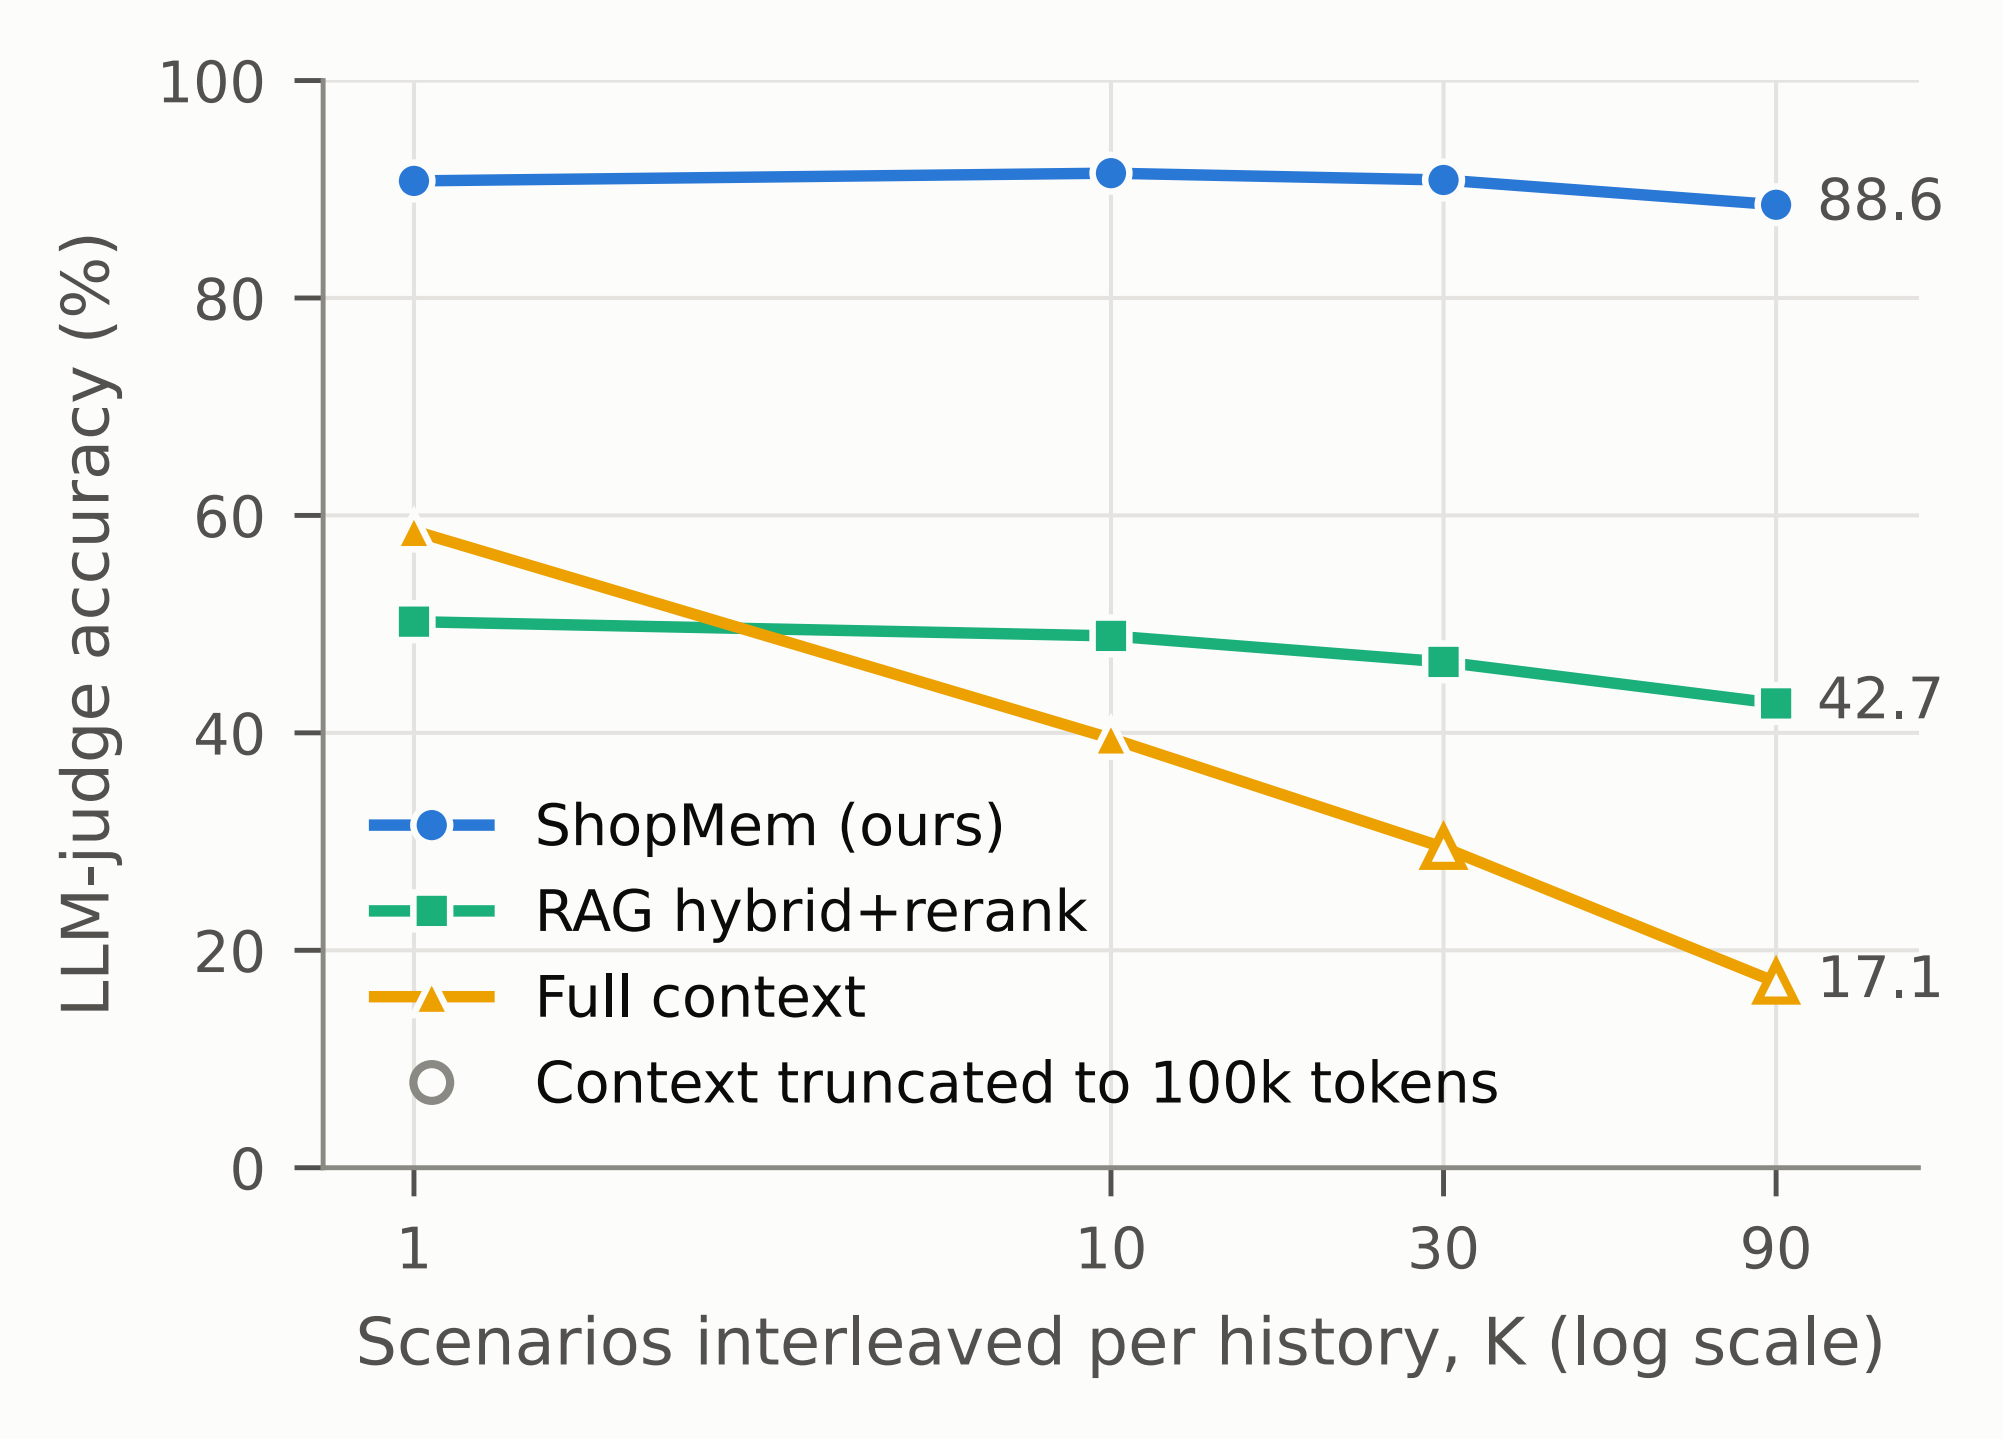
100
80
60
40
20
0
LLM-judge accuracy (%)
88.6
42.7
17.1
ShopMem (ours)
RAG hybrid+rerank
Full context
Context truncated to 100k tokens
1
10
30
90
Scenarios interleaved per history, K (log scale)