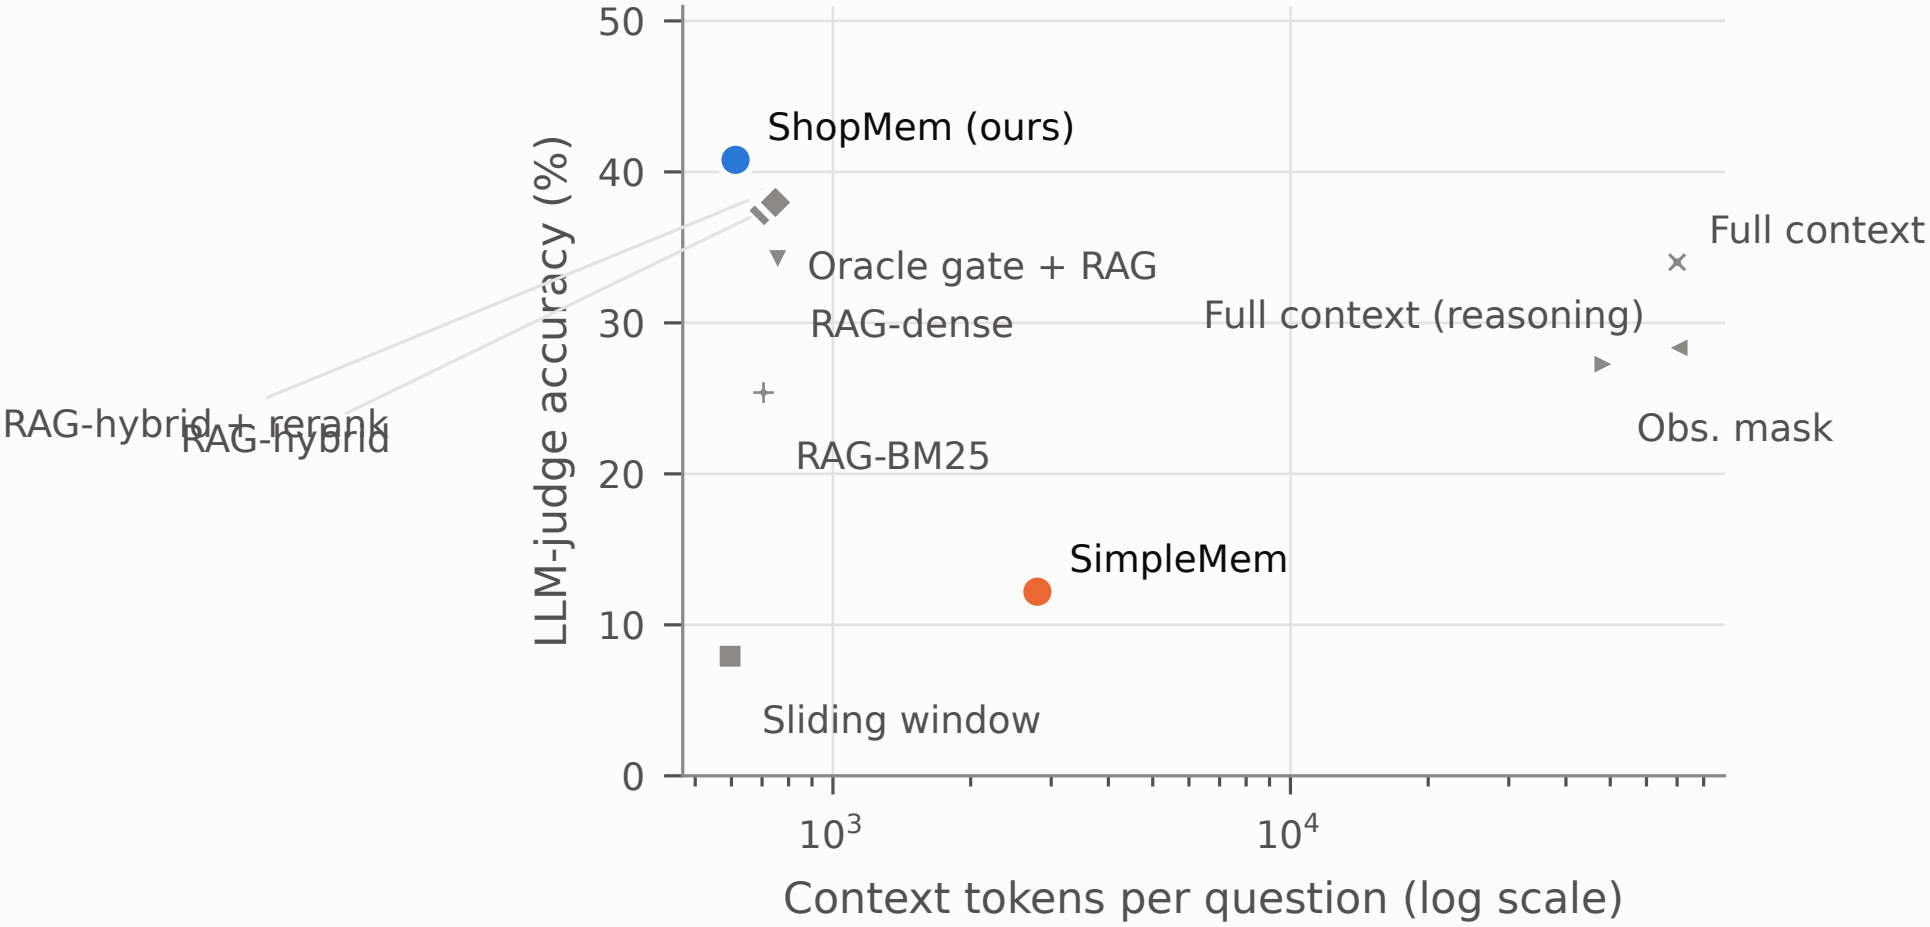
LLM-judge accuracy (%)
50
40
30
20
10
0
10^3
10^4
Context tokens per question (log scale)
ShopMem (ours)
Oracle gate + RAG
RAG-dense
RAG-BM25
RAG-hybrid + rerank
RAG-hybrid
Sliding window
SimpleMem
Full context
Full context (reasoning)
Obs. mask