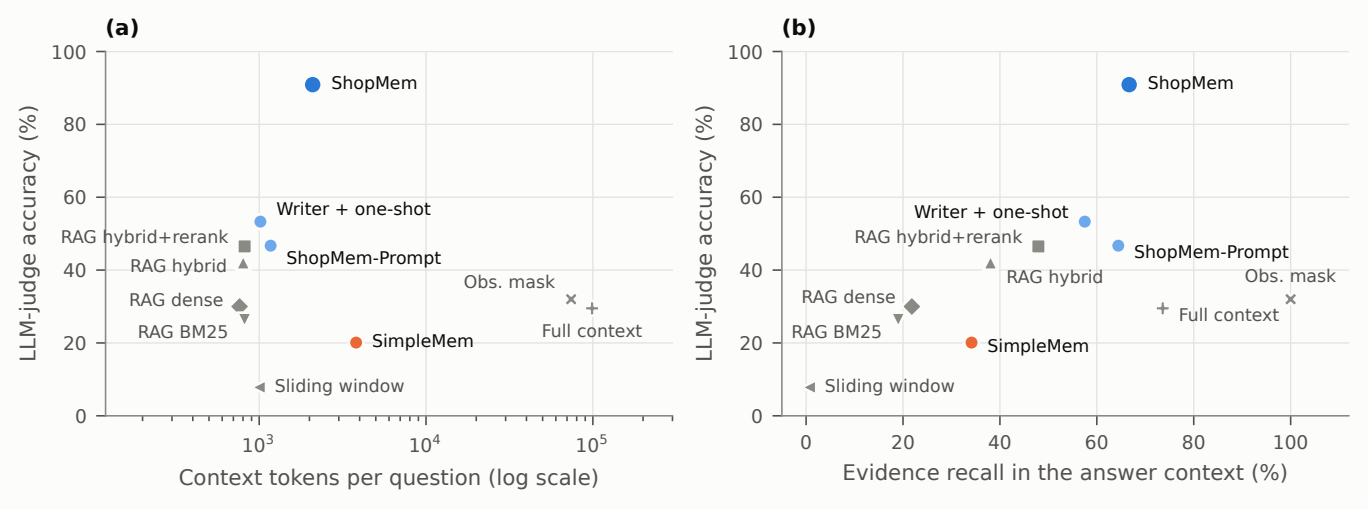
(a)
LLM-judge accuracy (%)
Context tokens per question (log scale)
ShopMem
Writer + one-shot
RAG hybrid+rerank
ShopMem-Prompt
RAG hybrid
Obs. mask
RAG dense
Full context
RAG BM25
SimpleMem
Sliding window
(b)
LLM-judge accuracy (%)
Evidence recall in the answer context (%)
ShopMem
Writer + one-shot
RAG hybrid+rerank
ShopMem-Prompt
RAG hybrid
Obs. mask
RAG dense
Full context
RAG BM25
SimpleMem
Sliding window