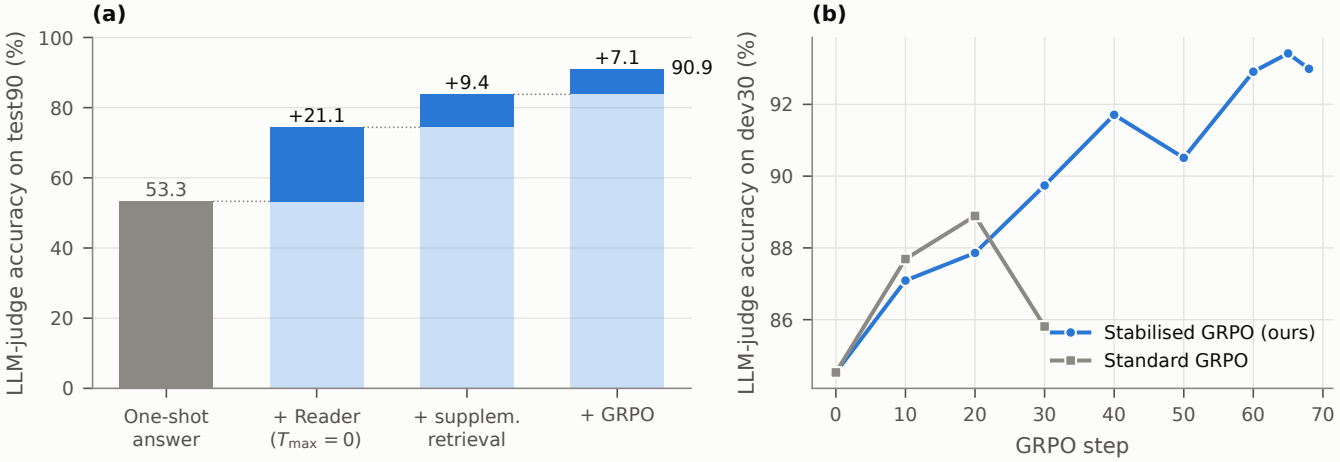
(a)
LLM-judge accuracy on test90 (%)
100
80
60
40
20
0
53.3
+21.1
+9.4
+7.1
90.9
One-shot answer
+ Reader ($T_{max} = 0$)
+ supplem. retrieval
+ GRPO
(b)
LLM-judge accuracy on dev30 (%)
92
90
88
86
Stabilised GRPO (ours)
Standard GRPO
0
10
20
30
40
50
60
70
GRPO step