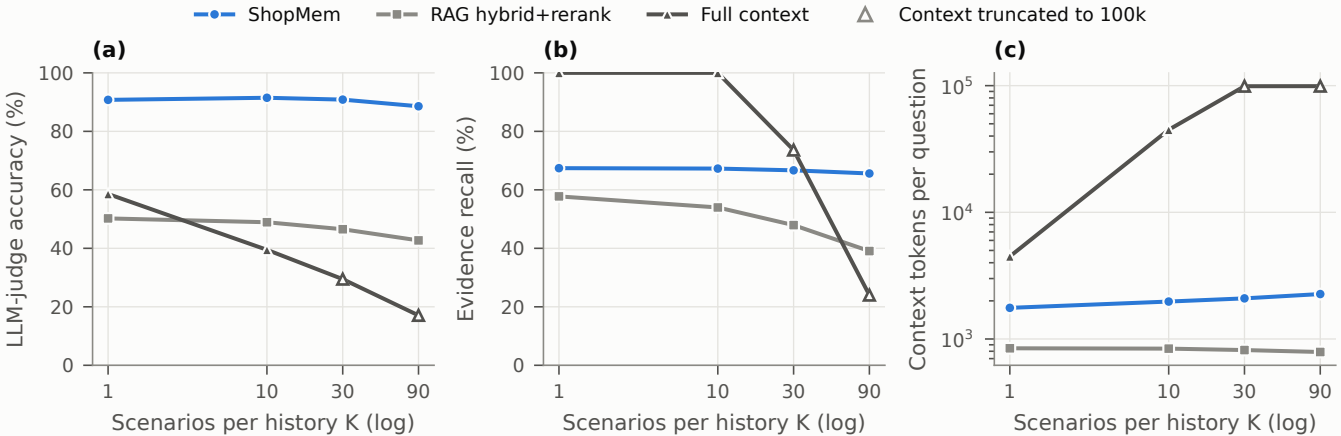
ShopMem
RAG hybrid+rerank
Full context
Context truncated to 100k
(a)
LLM-judge accuracy (%)
Scenarios per history K (log)
(b)
Evidence recall (%)
Scenarios per history K (log)
(c)
Context tokens per question
Scenarios per history K (log)