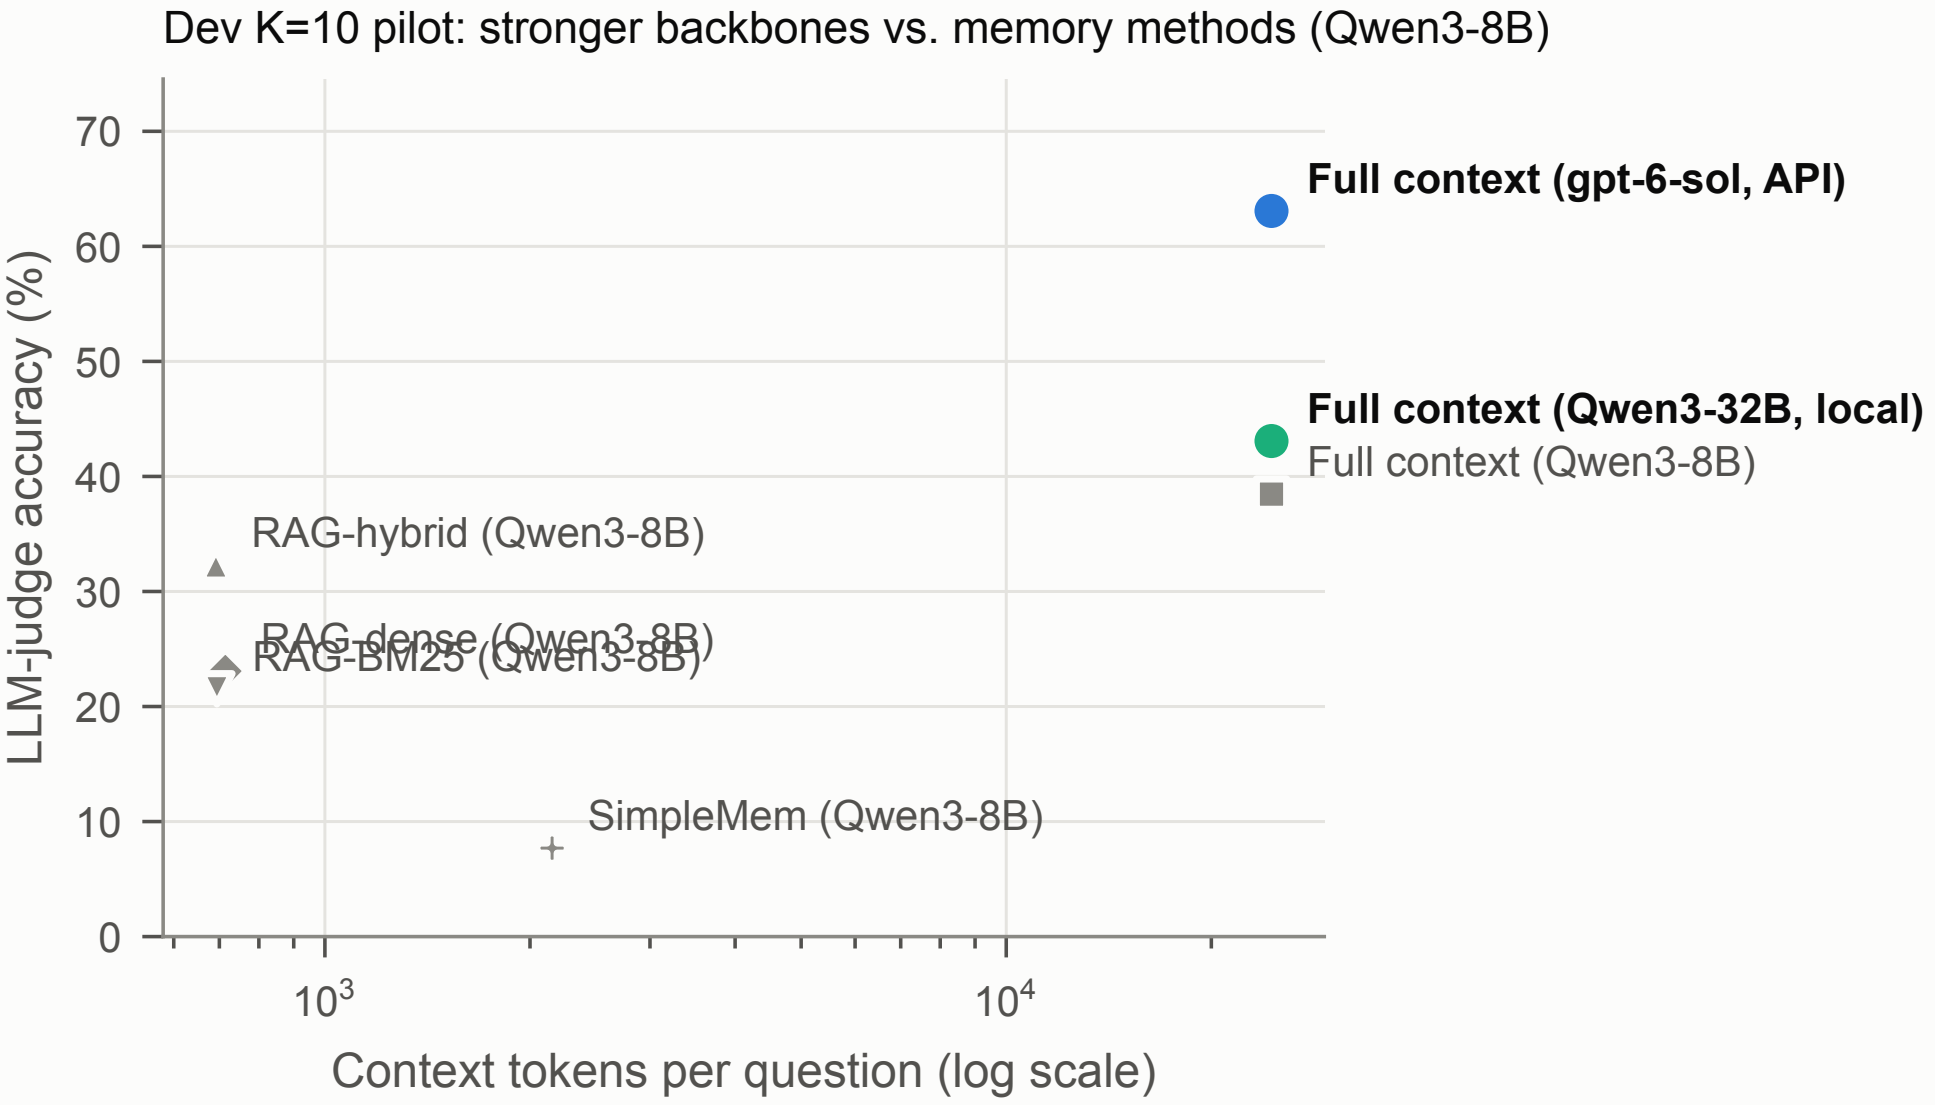
Dev K=10 pilot: stronger backbones vs. memory methods (Qwen3-8B)
LLM-judge accuracy (%)
70
60
50
40
30
20
10
0
10^3
10^4
Context tokens per question (log scale)
Full context (gpt-6-sol, API)
Full context (Qwen3-32B, local)
Full context (Qwen3-8B)
RAG-hybrid (Qwen3-8B)
RAG-dense (Qwen3-8B)
RAG-BM25 (Qwen3-8B)
SimpleMem (Qwen3-8B)